

Analysis of the common recommendation systems with the common frameworks: Spark and Flink

Final Report for the BigData project

Mirko Morandi
University of Trento
176043
mirko.morandi@unitn.it

Zhiheng Xu
University of Trento
174222
zhiheng.xu@unin.it

ABSTRACT

In this paper we provide an analysis of the state of the art of recommendation systems.

Collaborative Filtering is the current buzzword in the world of recommendations, came to notoriety after the Netflix Prize challenge. In this paper we aim to analyze the current implementations of two different algorithms used for Collaborative Filtering: **Alternating Least Squares** (ALS) and **Stochastic Gradient Descent** (SGD) in combination with the common frameworks **Spark**.

Keywords

Spark; CF; Collaborative Filtering; ALS; SGD; Scala

1. INTRODUCTION

Recommender systems are now trending due to the overwhelming availability of data. These systems have the ability to discover hidden relationships between users and items, and use these patterns to improve the user's taste prediction. Researchers discovered a "neighbourhood" of users with a similar taste which can be revealed by their previous actions: both implicit and explicit. *Collaborative filtering* is by far the most common approach adapted also by some of the biggest companies in the IT sector such as: **Amazon**, **Facebook** and **Netflix**. Amazon only recommendations accounts for the 35 % of **all sales**, a massive number if considered in proportion to their sales [5]. Although it's massive presence in the market, *CF* is not the only approach available for a recommendation system, but it is actually the successor of *Content-Based filtering*. The latter aims to profile a user searching the correlation with the item's peculiarity. By item peculiarity we refer to its implicit and explicit characteristics, for example a song's genre, subgenre, writer, composer, year of composition, beats per second etc. The problem with this approach lays in the difficult of retrieving all the necessary information, which sometimes are not even

available or disloable. Furthermore with the raise of the Big Data paradigm some frameworks started to grow from the academic world to the Apache Foundation: **Flink** and **Spark**. Those frameworks can be seen an extension of the Hadoop ecosystem, and both of them have their own pros and cons. In our paper we will use and briefly introduce only the latter one.

The paper is structured as follow: the description in details of what collaborative filtering is, pro and contro and alternatives. Then an extensive description of the application of Matrix Factorization to collaborative filtering, and the considered algorithms: *Alternating Least Square* and *Stochastic Gradient Descent*. Last part is our implementation with failures, successes and results.

2. COLLABORATIVE FILTERING

2.1 Collaborative Filtering Approaches

Collaborative Filtering can be subsequently defined in two different approaches:

2.1.1 Memory-based Collaborative Filtering

In *memory-based CF* uses user ratings to compute similarity between user and items and subsequently make a recommendation. Usually this approach involves "neighboring" algorithms such as **K-Nearest Neighbours** to build relationships between users. The similarity between two users is calculated using the **cosine similarity**.

Cosine similarity is a measure of similarity between two vectors of an inner product space that measures the cosine of the angle between them [2].

$$\cos(\mathbf{t}, \mathbf{e}) = \frac{\mathbf{t} \cdot \mathbf{e}}{\|\mathbf{t}\| \|\mathbf{e}\|} = \frac{\sum_{i=1}^n \mathbf{t}_i \mathbf{e}_i}{\sqrt{\sum_{i=1}^n (\mathbf{t}_i)^2} \sqrt{\sum_{i=1}^n (\mathbf{e}_i)^2}} \quad (1)$$

The recommendation is made by finding the top K similar users and aggregate their user-item matrices to find the appropriate recommendation. The typical problem of this approach is the difficult with scaling when the data gets bigger. Due to the *Big Data* paradigm expansion this approach has been deprecated favoring the following approach.

2.1.2 Model-based Collaborative Filtering

The most common approach to *CF* is through the factorization of a very big and sparse matrix. For example during the Netflix Prize at the participants were given a matrix of 8.5 billions of ratings, of which only 100 millions were non

zero values [3]. *Model-based CF* uses machine-learning and data mining algorithms to uncover the latent factor model between users and items to predict the missing ratings. **Latent factor models** are hidden relationships between users and items hardly discoverable in the original data; usually they may for example denote the quantity of action in a movie or the complexity of the characters. These vectors are then used to create the missing values in the user-items matrix.

Furthermore the model-based collaborative filtering can be expanded in two distinct sections: *user and item based collaborative filtering* depending on the priority given to the prediction.

2.2 Related problems

2.2.1 Cold start problem

Due to the nature of CF, the system needs a huge amount of data in order to produce a reliable prediction. But what happens if our system hasn't collected any or not enough information yet? This problem is called **cold start** and can be tackled with some advanced machine learning solutions called *active learning*.

2.2.2 Shilling Attacks

CF can be exploited to perturbate its prediction system with a technique called *shilling attacks*. This happens when the input system (e.g. ratings) are given in the correct way trying to alter the recommendations to the one favoured by the attacker. It has also been noticed that these kind of attacks affect more user-based CF algorithms instead of item-based.[10]

2.2.3 Sparsity

In the era of the *web 2.0* the matrices who compose the datasets are usually very sparse due to the typical proportion for which $nUsers \ll nItems$. As said previously the matrix which was given had only 100M ratings out of 8.5 billions of records. This problem is solved using matrix factorization.

3. MATRIX FACTORIZATION

During the last decade a huge effort has been applied to solve the problem of big datasets with an incredible amount of data. Let denote $R = r_{i,j}$ denote a user-movie where each element $r_{i,j}$ represents a rating given by a user i to an item j with a value from 0 to 5; where 0 means non rated and 1 to 5 a rating ranging relatively from *very poor* to *awesome*. Let also define m the number of users and n the number of movies in the system. The problem of *recommender systems* is to predict the missing values of R using the known ratings.

The concept behind *matrix factorization* is to find two matrices V, W relatively $m \times p, n \times p$ which product can approximate a much bigger matrix R of dimensions $m \times n$. $R \approx V * W^T$


The process consists in a low-rank approximation of the user-item matrix using for both users and items some feature vectors which are used to model the prediction with a inner vector product of the selected user and item.

Ideally $r_{i,j}$ should correspond to the predicted rating but in reality the value will be different. Therefore we need a loss function to determine the difference between **real value**

Table 1: Example of a sparse user-item ratings matrix

items/users	U1	U2	U3	U4	U5
I1	5	3	-	1	3
I2	4	-	2	1	-
I3	2	2	-	5	-
I4	4	3	-	4	2
I5	-	5	5	4	5

Figure 1: Matrix factorization



and **predicted value**. In our paper we are going to use the **Root Mean Square Error**,

$$\sqrt{\frac{1}{N} \sum_{i,j}^N (pr_{i,j} - r_{i,j})^2} \quad (2)$$

Where respectively $pr_{i,j}$ is the prediction for the rating given by the user i for an item j and $r_{i,j}$ is the real rating for the former user and item. The prediction is obtained with the dot product of the user-item vectors, (u_i, v_j)

$$prediction_{i,j} = \bar{u}_i * \bar{v}_j \quad (3)$$

The low rank approximation problem is formulated as follows to learn the factor vectors (u_i, v_j) [1]

$$(u_i, v_j) = \min_{u,v} \sum_{(i,j) \in K} (pr_{i,j} - r_{i,j})^2 \quad (4)$$

However solving the approximation with the former approach leads to the overfitting of data. Overfitting means that our model is learning through the noise in the training data instead of the underlying relations for which we expect to work. Although the performance will continue improving with the training data it will get worse on unseen data. This usually happens when we are using a too complex model or we have too many free parameters, and the latter is our case. The solution is to apply regularization penalizing the magnitude of the feature vectors proportionally with a constant factor λ obtaining the following and well known formula for low-rank approximation:

$$(u_i, v_j) = \min_{u,v} \sum_{(i,j) \in K} (pr_{i,j} - r_{i,j})^2 + \lambda(u_i^2 + v_j^2) \quad (5)$$

3.1 Approaches

There are mainly two approaches to minimize equation(5) and are known as:

Alternating Least Square and Stochastic Gradient Descent.

3.2 Stochastic Gradient Descent

Stochastic Gradient Descent is a gradient descent optimization used to minimize a loss function, in our case equation (5). The goal of *Stochastic Gradient Descent* is to find a value $\theta^* \in R^k (k \geq 1)$ that minimizes a given loss $L(\theta)$. The algorithm makes use of noisy observations $\bar{L}'(\theta)$ of $L'(\theta)$, the function's gradient with respect to θ . Starting with some initial value θ_0 , SGD refines the parameter value by iterating the stochastic difference equation [7]

$$\theta_{n+1} = \theta_n - \epsilon_n \bar{L}'(\theta_n) \quad (6)$$

where n denotes the step number and ϵ_n is a sequence of decreasing step size.

3.2.1 SGD with Matrix Factorization

In matrix factorization we set $\theta = (V, W)$ and the loss function can be written as $Q(w) = \sum Q(u, i)$. Where w is an approximation obtained iterating through every single sample until a minimum is reached with the following step: $w = w - \alpha * d_w^d(Q_i(w))$ In our case $Q(u, i)$ is represented by equation(5), and is the equation we want to minimize. The iterative version of the algorithm can be sketched this way:[4]

Algorithm 1 Matrix Factorization with SGD

Require:

R is the user-item matrix,
V and W are the factor vectors initialized with values from 0.0 to 1.0,
 α is the learning rate,
 λ the magnitude reduction,
K is the number of iterations
F is the number of features

```

1: shuffle(Ratings)
2: for i to K do
3:   for user u and item i with a rating  $r \in R[u, i]$  do
4:      $predictedRating = V[u] * W[i]^t$ 
5:      $error = R[u, i] - predictedRating$ 
6:     for each feature  $f \in F$  do
7:        $V[u, f] = V[u, f] + \alpha * (error * W[i, f] - \lambda * V[u, f])$ 
8:        $W[i, f] = W[i, f] + \alpha * (error * V[u, f] - \lambda * W[i, f])$ 
9:     end for
10:   end for
11: end for
```

Although this algorithm is used to optimize equation(5) it has some scalability issues. These problems are insights of its iterative nature for which all the individual steps depend on each other and make it hard (but not impossible) to distribute it. In this paper we will propose a functional implementation of this algorithm, and try to optimize it with a distributed version.

3.3 Alternating Least Square

Because both u_i and v_j are unknowns, Equation 5 is not convex. However, if we fix one of the unknowns, the optimization problem becomes quadratic and can be solved

optimally. ALS is the technique that rotates between fixing the u_i 's and fixing the v_j 's. When all u_i 's are fixed, the system recomputes the v_j 's by solving a least-squares problem, and vice versa. This ensures that each step decreases Equation 5 until convergence. In this section, the cost function is defined as,

$$Q(U, V) = \sum_{(i,j) \in K} (pr_{i,j} - r_{i,j})^2 + \lambda(u_i^2 + v_j^2) \quad (7)$$

The detailed description is as following,

Algorithm 2 Matrix Factorization with ALS

Require:

R is the user-item matrix,
U and V are the factor matrixes initialized with values from 0.0 to 1.0,
 α is the learning rate,
 λ the magnitude reduction,
K is the number of iterations
rmse is expected RMSE
 r_i is the i th row of R, r_j is the j th column of R

```

for k to K or RMSE > rmse do
2:   fix V, calculate the partial derivative for  $u_i$ 
   make it equal to 0, then have
4:    $u_i = (V^T V + \lambda I)^{-1} V^T r_i$ 
   update all the  $u_i$ 
6:   then fix U, calculate the partial derivative for  $v_j$ 
   make it equal to 0, then have
8:    $v_j = (U^T U + \lambda I)^{-1} U^T r_j$ 
   update all the  $v_j$ 
10: end for
```

One obvious advantage of ALS is that the system can easily use parallelization. In ALS, the system computes each u_i independently of the other item factors and computes each v_j independently of the other user factors. It means that the algorithm can update all the u_i and v_j in parallel. This gives rise to potentially massive parallelization of the algorithm and improve the training efficiency.

3.3.1 Weighted Alternating Least Square

Weighted alternating least squares(W-ALS) is an improved algorithm of ALS. W-ALS aims to train the recommendation system with implicated dataset. In some situation, users don't rate items exactly with values but through their behaviour, which is called implicated ratings. An example of implicit ratings is that when people are watching TV program, they don't actually rating, but the recommendation system can assign user's ratings basing their watching time on the program. W-ALS introduces different confidence levels for which items are preferred by the user changing the objective function to be minimized to:

$$(u_i, v_j) = \min_{u,v} \sum_{(i,j) \in K} c_{i,j} (p_{i,j} - r_{i,j})^2 + \lambda(u_i^2 + v_j^2) \quad (8)$$

Here $c_{i,j}$ measures the confidence of the predicted preference $p_{i,j}$. There are many ways to compute the confidence level, but it is out of the discussion in this paper.

3.4 Machine Learning and BigData

Machine learning is ideal for exploiting the opportunities hidden in big data. It delivers on the promise of extracting value from big and disparate data sources with far less reliance on human direction. It is data driven and runs at machine scale. It is well suited to the complexity of dealing with disparate data sources and the huge variety of variables and amounts of data involved. And unlike traditional analysis, machine learning thrives on growing datasets. The more data fed into a machine learning system, the more it can learn and apply the results to higher quality insights. Freed from the limitations of human scale thinking and analysis, machine learning is able to discover and display the patterns buried in the data.[8] Data by itself is useless if there is no way to extract information from it. That's why Big-Data goes always in couple with machine learning or data mining; these two different but linked branches of computer science aims to extract or predict relevant data from a huge amount of information. *Data Mining* is a field of computer science where algorithms tries to find patterns and correlations analyzing data; there are some subcategories in data mining known as *text mining* and *process mining*. Text mining is used to find correlations between words in a text file, the most famous example is a search engines; search engines needs to find the appropriate result matching the keywords analyzing the words in a webpage, the most common approach is **PageRank**. We can also use data mining to extract relevant information from our log files in order to detect errors, bottlenecks or unusual behaviours (attacks or viruses); this is called **process mining**.

Machine learning is another branch of computer science used to make predictions after "teaching" the machine how to learn from data. The typical approach is to use a dataset called *training data* (usually 80% of the dataset) to build a model. The training phase is where our algorithm makes predictions using the data and compute the error relative to the real value, and subsequently improve the predictions until convergence (e.g. when reaching a threshold for which the error is acceptable). After reaching the convergence threshold we have to test our data with unknown data (usually the remaining 20% of the dataset) and then compute the predictions on the latter. Typically the error is slightly higher with the test dataset (if we have the same error it means we have a wrong model), otherwise if the difference is significant it means that we are *overfitting* the model. Overfitting means that we are training our model wrongly and it does not learn from the hidden correlations between values but instead from the noise relate to that.

4. TECHNOLOGIES AND CHALLENGES

Up to now we have described what is the state of the art in this branch of the reserarch. Our challange in this paper will be to develop, where needed, and compare the two most common approaches to matrix factorization using the *Spark* framework. There's already an implemented version of the *Alternating Least Square* with scala, which we'll use to do our tests. Our goal will be to implement from scratch the *Stochastic Gradient Descent* starting from the naive iterative version and optimize it aiming to reach a distributed version. We have also choose to use **Scala**, a functional programming language to develop the SGD factorization; functional programming languages seems to be better involved in the bigdata paradigm due to their ability to facilitate parallelization. Last but not least, we have

used the various movielens datasets with different sizes to test our examples.

4.1 Scala

Scala, acronym of "Scalable Language", is an object-oriented functional programming language used for a large variety of tasks: from scripting to large mission critical systems (e.g. Twitter and Linkedin). Functional programming languages are well known in the bigdata environment due to one of their peculiarities: **The pure functions**. Pure functions are like mathematical functions, the return value depends only on the input value; meaning the function does not have any *side effect*. This is very important when dealing with large datasets because it's the peculiarity of functional languages that improves parallelization. The requirement for scaling horizontally is to have many independent tasks running in parallel, which goes perfectly with the latter peculiarity of functional languages. However Scala offers much more than pure functions, hereby follows some of the most important aspects of this programming language:

- Java Interopability: Scala compiles on a JVM, which means that basically we can use all the libraries written for Java
- Strong static typing: helps the developer to detect bugs and errors at compile time by design
- Type inference: the type of the variable doesn't have to be declared but it is inferred.
- Pattern Matching: pattern matching is a useful tool for guiding the code flow using patterns to match which branch to follow in the code.
- Non-alphanumeric identifiers: operators such + or - can be used as **add** or **subtract**, or there can be more user defined (see the dot product between vectors from breeze).
- Anonymous functions as objects
- The expression oriented paradigm: all parts of the code of a Scala program yields a result, which means they're all expressions and **not statements**. The difference is subtle but essential: statements does not yield a result, they just does something; e.g. `double gamma = 10;`. Meanwhile expressions are always evaluated, for example the if-else does always have to return a value, when this is not true for object oriented languages; e.g. `val gamma = if(condition) 3 add 5 else pow(2,2).`
- reassignable variables: not common in functional programming Scala supports the riassignment of variables with the keyword **var** instead of **val** (although its use is recommended only if strictly needed).
- tail recursion: computing a recursion as last operation of the function is called tail recursion; it's known for improving performance and memory usage.

4.2 Spark

Spark is an open-source processing engine for distributed data processing and data analytics. Spark was born as a university project to solve some performance problems with

Table 2: MovieLens datasets

-	Ratings	Users	Movies
1 Million dataset	1.000.209	6.040	3.900
Full dataset	22.775.244	246.829	33.869

Table 3: Book-Crossing datasets

-	Ratings	Users	Books
Full dataset	1.149.780	278.858	271.379

Hadoop on a large scale, and rapidly became one of the most widely adopted frameworks for data analysis.

Spark offers an ecosystem with the support of the most innovative tools such as *Streaming*, machine learning libraries for big datasets (and also for streaming, which is not available for the competitor *Flink*), dataframes libraries and many other features.

Although Spark has on it's own ecosystem, one of the reasons we choose to use it instead of it's competitor Flink is due to the number of libraries available on the internet and the wide support given by the community. The latter is the biggest problem with Flink, which has basically no online support due to it's relative youth, and made us choose Spark over it.

4.3 Libraries

We used a variety of different softwares in order to take care of some crucial aspects of our project such as performance and deployability.

Spark 1.6.0

Scala 2.10.6

Sbt 13.8 for managing dependencies and packaging the project

Spark MLlib 2.10 for machine learning tasks

Breeze 11.2 for scala, for machine learning optimization

Mahout 0.9 for other machine learning optimization

4.4 MovieLens Dataset Structure

In our project we used mostly two datasets from *MovieLens*: the 1 million ratings and the full dataset, as shown in Table 2.

4.5 The Book-Crossing dataset

In order to compare the algorithms performances we have also added another dataset for collaborative filtering (one of those who are free). It's a dataset of book ratings ranging from 0 to 10, where 0 means not rated.

4.6 The machines

During our test we used different machines for testing and for production (in other words real external server). See Table 3 for more details

5. IMPLEMENTATION

The first step of our research was to develop a "naive" version of SGD using the various iterations with no synchronization. We also have decided to start using the smallest dataset from movielens to understand how long it will take

Table 4: Computer

OS	#CPU	#Threads	RAM
Mac OS X	2	4	8 GB
Ubuntu 14.04 remote	4	8	8 GB

to complete. The decision was wise because on the first run took around 8.7 hours to complete with a dataset of 1 million of ratings. The overall problem was already introduced in the previous chapters of this paper and it's the nature of this algorithm to not scale well with big datasets. Moreover if we consider our example and using $nIterations = 25$, $nFeatures = 20$ with $nUsers = 6000$ and $nMovies = 4000$ we would have to compute $nIterations * nUsers$ multiplied by the ratedmovies the dot product of the user item factor vectors, and update the values for $nFeatures$ times. We can summarize the cost of our iteration as $O(nIterations * nUsers * ratedMovies * nFeatures)$. Since the number of users is typically the highest value this doesn't scale well with big datasets where the number of users reaches very high numbers and the number of computations becomes enormous. Although the problem of scaling is relative to the algorithm we found out that some commonly used data structures and operations we used had a high impact on the performance.

5.1 The algorithm

In our implementation we used a variant of the simple SGD called biased SGD which considers also some bias related to users and items. Biases are really important in our example, because ratings are personal driven decisions makes them differ from person to person. Let's say that the overall average is 3.6 (that's the real average in our dataset) and Start Wars is a movie above the average which means it tends to be rated 0.5 more than other films. On the other hand we have a critical user Bob, who tends to give lower than average ratings to movies around 0.3, in our case the starting point for the prediction would be $(3.6 + 0.5 - 0.3) = 3.8$. The predicted rating will be $predictedRating = \mu + userbias + itembias + vectorFactorProduct$

Epoch Loop The first step of the algorithm is to iterate a given number of epochs and update the users. This is the basic loop

User-Item Loop The user-item loop is executed each time inside the former loop. This is a nested loop which goes through all the user and for each of them all the items for which they do have a non-zero rating. For each of this couple of user-items we first predict the value for the item and compute the relative error as follows: $error = predictedValue - realValue$. The value is predicted with the dot product of the factor vectors associated to the relative user and item. Here we also update the biases related to both user and items with the following formula $bias = bias + alpha * (err - biasRegulator * preventOverfitting * bias)$. This allows us to adapt the user bias for any given rating; alpha is a float value, usually around 0.5 used to help the system's precision in the bias computation. PreventOverfitting is another float value used to decrease the chance of overfitting the data, in our example is al-

ways around 0.1. It has been already mentioned during this paper under the name of λ , used to reduce the impact magnitude of each rating.

Feature Loop The feature loop is an iteration for all the features associated to the user and item vector factor; since the number of features is the same for both user and item factors we can solve everything in a single loop.

5.2 The optimization game

In our first try we made a large use of lists and other structures which didn't perform very well on a large scale. Hereby there's a table of the common mistake we made during the first stage development and how we solved it.

Lists The first mistake we made was to use simple scala *Lists*. Although their simplicity of use they share a common problem: the cost of retrieving an element. The cost of getting an element from a List is $O(n)$, which means the bigger the array the more time it will take to get an element. Luckily Scala gives also other much better performing structures such as *Arrays* and *Dictionary*, whose cost of getting an element is respectively $O(C)$ and $O(eC)$. The first one performs slightly better on a large scale but they do both perform much better than Lists.

Array Although Arrays are good for retrieving and updating elements they're not so good, on a large scale, for the dot product of vectors. We tried to perform a dot product of two Arrays of 5 millions of random values and two *DenseVectors* from the *breeze* library and the result was astonishing: 4948115 vs 39138 microseconds. We ran the experiment multiple times and the magnitude of the difference was the same even using the *parallel collections* to compute the dot product.


RDD filtering We needed to retrieve at each step the real rating from the dataset and order to compute this operation we filtered directly the distributed Dataset. Even though the dataset is distributed it had a very high cost of information retrieval, a value around 132080 microseconds. We have then decided to cache all the ratings using an *md5 digest* computed using the *userID* a space and the *movieID* to have a unique value (e.g. "2293 345"). This operations has a non negligible cost but after the initial setup the improvement was noticeable: 48 microseconds compared to the previous 132080.

6. RESULTS

The results from our tests proved the correctness of our implementation of SGD although the performance are slightly worse if compared to the ALS implementation.

1. The first test was run on a local machine with 25 iterations and 15 features. In picture(2) we can see the slow convergence with the typical curve of SGD with a very high peak in the beginning and then a fall until convergence to the real rating, 5.0 in our case. The overall iteration took around 45 minutes to complete with an acceptable error: $rmse = 1.09$ (using a reduced version of the dataset)

Figure 2: SGD factorization convergence



2. The second test was a run on the remote server using the extended version of the dataset with 50 iterations and more features, 30. The algorithm with no optimization took 8.7 hours to complete with an $rmse = 1.078$. The result was good but the running time was out of any acceptable range
3. The third test was run again on the remote server, this time with a highly optimized version of the algorithm. This time the number of iterations was set to 200 with 30 features. The running time was 1.78 hours with an $rmse = 0.92$. Compared to the previous test, the running time would be 44 minutes, which is a big improvement.
4. The fourth test was set up on a local machine with an optimized version of SGD and the 1M dataset, with 25 iterations and 20 features. The running time was an astonishing 8.57 minutes, a very good improvement compared to the previous ones ($rmse = 0.86$).
5. The last test was set up with the parallel version of the algorithm with the same parameters as the previous. The algorithm took 45 minutes to complete the whole process which is an improvement of the 400% compared to the previous version.

6.1 Playing with the RMSE

The most important factor and indicator of the whole algorithm is the *root mean square error*. The $rmse$ indicates the "correctness" of the algorithm, in our case a good value would be from 0.8 to 1.1 at maximum. The typical value for this dataset is around 0.85 [6].

Although the algorithm was correct in our first batch of test the $rmse$ was around 1.29, which is not the best value we could get. We tried to fine tune some parameters, starting with the number of iterations and the number of features, the ones we considered the most relevant in the algorithm. However we were wrong, the error was not decreasing as expected but only of few decimal values. The key to our high value was in the *bias learning rate* or α in the pseudocode, which was too high and was influencing too much the prediction. Using a value around 0.01 we decreased the error of a factor of 32%, with a resulting $rmse = 0.88$ on the test

Figure 3: SGD RMSE comparison


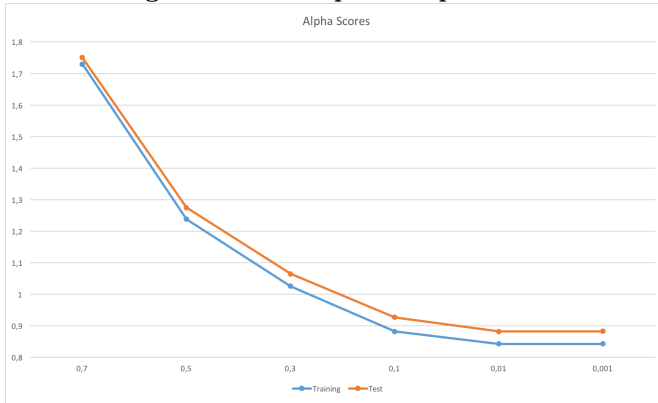


Figure 4: SGD alpha comparison



data and 0.84 on the training data, which is a good result (see figure 3).

6.1.1 Tuning alpha

As said previously we found out that the best α value (*bias learning rate* in our code), is approximately 0.01. In this paragraph we will analyze the algorithm behaviour with different values starting from 0.7 (the numbers above caused some integer overflows).

As can be seen in the figure, after a certain value the error stop decreasing and it stabilizes slowly to a certain threshold 0.88, and it becomes useless to lower the value.

6.2 Investigation over other parameters

6.2.1 The λ value

As said previously the λ value is used to reduce the magnitude of each rating preventing the overfitting of our dataset. We have run this test with the 1M MovieLens dataset on our local machine. The results highlighted a correlation between the speed of convergence to the "correct" error and the λ value. The higher and the slower will be the convergence to a better result.


6.2.2 The number of features

We thought that the number of features would be on of the most valuable parameters which could make the difference. However we have noticed that increasing or decreasing

Figure 5: SGD learning rate



Figure 6: SGD learning rate decay



their number (10 or 50) the difference was very subtle. We have tried then to raise significantly the number of features to **2000** to see how they would have impacted the overall performances. In this case we have used an α parameter of 0.5, to see if the number of features would have improved something. The results proved we were correct; the error was not improving as expected, the improvements were negligible (around 0.01) and the performance degraded a lot.

6.2.3 The learning rate

The learning rate value is the value from where to start the learning process. In all of our tests it has been set to 0.1, but we have decided to investigate its influence on the algorithm. The range of values is between 0.175 and 0.05 (higher values causes overflow). We discovered that the starting value influences the speed of which the algorithm starts to converge to an error.

In this test we have used on purpose a high α value (0.5) to test wheter the algorithm can improve.


6.2.4 Learning Rate Decay

The learning rate is a parameter that determines how much an updating step influences the current value of the weights. We have noticed that increasing this value led to a faster reduction of the error, although there are some spikes in the error dropping. The best value for it is around 0.9 - 1.0, the numbers above causes errors.

6.3 Comparison with the non-biased version

As said in the previous chapters there are many optimized versions of SGD, including the one used during the whole

Figure 7: SGD implementations



paper. We have considered to be of interest a comparison between the two versions, in a matter of running time and error.

The running time did not differ too much from it's improved version, though there is a substancial difference in the rmse calculation. The non-biased version starts from a higher value (3.7 compared to 1.12) and stops decreasing the error converging slowly to a higher error (1.06 compared to 0.86).

6.4 Profiling ALS

We give less space to this algorithm since it's behaviour and performance is well known, since is the most adopted for this task. However it is a good metric of comparison between our implementation of SGD.

We have conducted several test with respect to: error of predictions and performances. The basic assumption is that the test is run with $\lambda=0.1$, number of features=10, number of iterations=5 on the the smaller dataset presente in the table above (1M ratings)..

6.4.1 Number of Iterations

The first test is to compare the error on predictions of the algorithm related to the number of iterations. We could do at maximum 20 iterations before the machine break down with a stack overflow message. However the error was quite low and in the end converged around 0.86, as can be seen in **Figure 8**.

6.4.2 Number of Features

The second test is to compare the error on predictions of the algorithm related to the number of features. At the beginning, we tested with iteration=5 and with number of features respectively equals to 8, 10, 12 and 15, but we found that the number of features doesn't influence the RMSE a lot. Thus, we also tested with iteration=20, and we got the similar result in **Figure 9**. As result, we could say the the number of features doesn't influence the RMSE that much if the number of features is around 10.

6.4.3 Value of λ

The third test is to compare the error on predictions of the

Figure 8: ASL RMSE with different iterations




Figure 9: ASL RMSE with different number of features




Figure 10: ASL RMSE with different λ




Table 5: Running times

Matrix	LocalALS	SparkALS	ParalleledALS
3.900 * 6.040	about 2792s	about 5400s	about 52s

algorithm related to the value of lambda. We tested with λ respectively equals to 0.08, 0.1, 0.3 and 0.5. As we can see from **Figure 10**, the bigger λ we have, the worse RMSE we got.

6.4.4 Full Dataset

The fourth test was a performance test on the complete dataset (22M ratings). The algorithm behaved very well, with an average running time of 4.8 minutes. Furthermore we have noticed a reduction of the RMSE (10 iterations) to the value 0.81, which is the best one we had in all the tests so far.


6.4.5 ALS Versions

We have also conducted a test to investigate the different running times between the three different versions of ALS : the LocalALS without using Spark API, the SparkALS partial using Spark API which is partially paralleled, and ParalleledALS fully using Spark API which is fully paralleled. Running on a Matrix with 3.900 movies, 6.040 users, the training has 5 iteration. For the SparkALS, it is run in 2 slices as when we try to run with 5 slice we got out of Memory error. The running times are shown in **Table 5**. It is interesting that the partial paralleled version has the worst performance, probably because it is just partial paralleled without optimization. On the other hand, ParalleledALS sharply improves the running performance with 5 slice(blocks), which has the best performance because the Spark API is already optimized.

6.4.6 Number of Blocks

As the ParalleledALS has the best running time performance, we run a test to see how number of blocks influences the running time. According to **Figure 11**, instead of the performance improved as the number of blocks increases, the ParalleledALS got the best running time(about 52s) when the number of blocks is equal to 5. We supposed the reason is that when blocks increases, the shuffle cost also increases and the shuffle cost overwhelms the improvement with more

Figure 11: ASL RMSE with different number of blocks



parallelization.

6.5 The parallel version

We analyzed the iterative version and in order to proceede in the parallelization we had to find some indipendent parts in our code. We made the assumption that all users may be processed indipendently even though the items will be shared. We solved this problem using concurrency on the items allowing only one of the parallel executions to have it. We have sliced the execution assigning theoretically a thread to each slide until the environment allows us to. In our case the maximum was 1500, but with some fine tuning of the heap and stack parameters one can allocate around 100 thousands of Java threads.[9]

Although the parallel version of our algorithm needs some locking system due to the concurrency on the items vector when reading and updating it there is a theoretical approach to a more distributed version. The key of this approach is to create a graph with all the interdependences between users and items, where the edges are the correlations. If we have more connected components, those will be fully independant from the others because there will be no concurrency issue between items, therefore we will be able to compute those on different machines and collect the results in the end. The only problem is relative to the possible number of connected components which is by nature of the problem very small, if not only one.

6.6 A different dataset

We have run some of the test on the Book-Crossing dataset to verify the behaviour of the two algorithms and then compare it with the former results.

6.6.1 SGD on Book-Crossing

The algorithm performed much faster compared to the *MovieLens* datasets, with an excellent rmse value. Although the number of users and books is much higher the algorithm performed much better. The reason behind this different performance can be found in the different sparsity of the matrices taken in consideration, that's the only reason we found to be acceptable. We have also noticed a consistent difference between the *rmse* we found compared to the oth-

Algorithm 3 Matrix Factorization with parallel SGD

Require:

R is the user-item matrix,
 V and W are the factor vectors initialized with values from 0.0 to 1.0,
 α is the learning rate,
 λ the magnitude reduction,
 K is the number of iterations
 F is the number of features
1: shuffle(Ratings)
2: **for** i to K **do**
3: run a thread for each user until the maximum number or the limit is reached
4: **for** each user u **do**
5: **for** each item i rated by u **do**
6: lock(item i)
7: $predictedRating = V[u] * W[i]^t$
8: $error = R[u, i] - predictedRating$
9: **for** each feature $f \in F$ **do**
10: $V[u, f] = V[u, f] + \alpha * (error * W[i, f] - \lambda * V[u, f])$
11: $W[i, f] = W[i, f] + \alpha * (error * V[u, f] - \lambda * W[i, f])$
12: **end for**
13: unlock(item i)
14: **end for**
15: **end for**
16: wait for all threads to join
17: **end for**

Table 6: Serial comparison: RMSE and Running

Dataset	Training	Test	Running time
MovieLens 1M	0.82826	0.868157	8.57m
Book-Crossing	0.438772	0.435893	4.278m

ers in the literature, with a consistence difference making us doubting about the correctness of our predictions. We checked some of the predictions on the test set and they were actually very close to the real value, though we have no clue of this difference between our result and the one gained by others.[1]

6.7 ALS on Book-Crossing

We were surprised to notice such a difference of results with the letter experiment. With this dataset the algorithm performed not too much better than the Serial version of SGD (4.27 minutes compared to 3.51 for this experiment). However, it's not the running time what surprised the most but the rmse difference, we had an average of 3.9 against the 0.4 of the previous experiment.

The only explanation we have found to the problem is in the number of ratings, which does not allow our algorithm to start working properly. Furthermore in the literature there are very similar values of rmse for this dataset and ALS. [1]

Table 7: Parallel comparison: RMSE and Running

Dataset	Training	Test	Running time
MovieLens 1M	0.828263	0.868157	4.37m
Book-Crossing	0.439313	0.436178	3.31m

Table 8: Comparison: RMSE and Running

Dataset	Training	Test	Running time
MovieLens 1M	0.759313	0.86346	0.37m
Book-Crossing	2.088756	3.836608	3.51m

7. CONCLUSIONS

Although the already existing implementations of collaborative filtering algorithms (ALS) have excellent running times we have developed a starting point for an SGD implementation in Scala which performances are of the same magnitude, minutes and not hours. This means optimization and parallelization is key to success, and there is always space for optimization, not always for parallelization though. It has been interesting to discover how the appropriate data structures can make such a big difference when using enormous quantities of data. What is used in the "classical programming" is never the same when applied thousands or millions of times, even the smallest operation can make the difference. Another very important aspect we learned to take care of is **indipendency**; indipendency of operations can make the difference between serial execution and parallel, and serious performance improvements. Being able to parallelize parts of the algorithm can reduce drastically the running time, both using more threads or machine clusters. However due to its definition, every SGD step depends to the next one and make it hard to be executed in parallel. This has required from us a some study and investigation on which parts can be made independant and which others can not, resulting with a *item-locked* implementation. Although the performances of the algorithm we have gained better results than the ones found in other works. There is a difference of respectively 0.16 and almost 2 with our *rmse* and the one resulting in [1].

8. FUTURE WORK

SGD is very promising and the research in this branch (collaborative filtering) is going faster than ever to adapt itself to the increasing size of datasets. The future work will be to have a fully parallelizable version of SGD, for which there are already some proposals of very complex and not naive algorithms.

9. REFERENCES

- [1] C. R. Aberger. Recommender: An analysis of collaborative filtering techniques.
- [2] W. C. foundation. Cosine similarity. https://en.wikipedia.org/wiki/Cosine_similarity.
- [3] S. Funk. Netflix update: Try this at home. <http://sifter.org/~simon/journal/20061211.html>.
- [4] A. Mahout. Introduction to matrix factorization for recommendation mining. <https://mahout.apache.org/users/recommender/matrix-factorization.html>.
- [5] M. Marshall. Aggregate knowledge raises 5m from kleiner, on a roll. <http://venturebeat.com/2006/12/10/aggregate-knowledge-raises-5m-from-kleiner-on-a-roll>.
- [6] Mymedialite. Mymedialite: Example experiments. <http://mymedialite.net/examples/datasets.html>.

- [7] P. J. H. Rainer Gemulla, Erik Nijkamp and Y. Sismanis. Large-scale matrix factorization with distributed stochastic gradient descent. *KDD 11 Proceedings of the 17th ACM SIGKDD international conference on Knowledge discovery and data mining*, 2011.
- [8] Skytree. Why do machine learning on big data? <http://www.skytree.net/machine-learning/why-do-machine-learning-big-data/>.
- [9] Stackoverflow. How many threads can a java vm support? <http://stackoverflow.com/questions/763579/how-many-threads-can-a-java-vm-support>.
- [10] X. Su and T. M. Khoshgoftaar. A survey of collaborative filtering techniques. page 19. Department of Computer Science and Engineering, Florida Atlantic University, 777 Glades Road, Boca Raton, FL 33431, USA, August 2009.

APPENDIX